

P5

Final Project: File System

1. Purpose

In this project, you will build a Unix-like file system on the *ThreadOS*. Through the use of the file system, user programs will now be able to access persistent data on disk by way of stream-oriented files rather than the more painful direct access to disk blocks with `rawread()` and `rawwrite()`.

2. Collaboration

Given the size of this project, you are encouraged to form a team of **two or three** students. In working with your partner(s), the workload should be divided equally across all members. This project requires much more time to be completed than the previous four assignments. Schedule and pace your design and coding work accordingly. If you are very confident of working independently you may do so.

3. What Files to Use

Use *ThreadOS*' original files, (i.e., the original classes) except for *Kernel.java*. For *Kernel.java*, use the one you got in the assignment 4, (i.e., *Kernel_org.java*). For simplicity, you should just re-download all readable files from the *ThreadOS* directory to get started.

4. Interface

The file system should provide user threads with the system calls that will allow them to format, to open, to read from, to write to, to update the seek pointer of, to close, to delete, and to get the size of their files.

For simplicity, the file system being created will consist of a single level. The "/" root directory is predefined by the file system and permanently available for user threads to store their files. No other directories are provided by the system and created by users dynamically.

Each user thread needs to keep track of all files it has opened. For this purpose, it should maintain a table of those open files in its TCB. This table is called a user *file descriptor table*. It has 32 entries. Whenever a thread opens a file, it must allocate to this file a new table entry, termed a *file descriptor*. Each file descriptor includes the file access mode and the reference to the corresponding file (structure) table entry. The file access mode indicates "read only", "write only", "read/write", or "append". The file

(structure) table is a system-maintained table shared among all user threads, each entry of which maintains the seek pointer and the inode number of a file. Depending on the access mode, the seek pointer is set to the first or the tail of the file, and keeps track of a next position to read from and to write to the file. It is entirely possible for one thread to open the same file many times, thus having several entries in the corresponding TCB's user file descriptor table. Although each of these user file descriptor table entries refer to a different file (structure) table entry with its own seek pointer, all of them eventually points to the same inode.

The file system you will implement must provide the following eight system calls.

1. **int SysLib.format(int files);** Formats the disk (*Disk.java's* data contents). The parameter *files* specifies the maximum number of files to be created (the number of inodes to be allocated) in your file system. The return value is 0 on success, otherwise -1.
2. **int fd = SysLib.open(String fileName, String mode);** Opens the file specified by the *fileName* string in the given *mode* (where "r" = read only, "w" = write only, "w+" = read/write, "a" = append). The call allocates a new file descriptor, *fd* to this file. The file is created if it does not exist in the mode "w", "w+" or "a". *SysLib.open* must return a negative number as an error value if the file does not exist in the mode "r". Note that the file descriptors 0, 1, and 2 are reserved as the standard input, output, and error, and therefore a newly opened file must receive a new descriptor numbered in the range between 3 and 31. If the calling thread's user file descriptor table is full, *SysLib.open* should return an error value. The seek pointer is initialized to zero in the mode "r", "w", and "w+", whereas initialized at the end of the file in the mode "a".
3. **int read(int fd, byte buffer[]);** Reads up to *buffer.length* bytes from the file indicated by *fd*, starting at the position currently pointed to by the seek pointer. If bytes remaining between the current seek pointer and the end of file are less than *buffer.length*, *SysLib.read* reads as many bytes as possible, putting them into the beginning of buffer. It increments the seek pointer by the number of bytes to have been read. The return value is the number of bytes that have been read, or a negative value upon an error.
4. **int write(int fd, byte buffer[]);** Writes the contents of *buffer* to the file indicated by *fd*, starting at the position indicated by the seek pointer. The operation may overwrite existing data in the file and/or append to the end of the file. *SysLib.write* increments the seek pointer by the number of bytes to have been written. The return value is the number of bytes that have been written, or a negative value upon an error.
5. **int seek(int fd, int offset, int whence);** Updates the seek pointer corresponding to *fd* as follows:
 - If *whence* is SEEK_SET (= 0), the file's seek pointer is set to *offset* bytes from the beginning of the file
 - If *whence* is SEEK_CUR (= 1), the file's seek pointer is set to its current value plus the *offset*. The *offset* can be positive or negative.
 - If *whence* is SEEK_END (= 2), the file's seek pointer is set to the size of the file plus the *offset*. The *offset* can be positive or negative.

If the user attempts to set the seek pointer to a negative number you must clamp it to zero. If the user attempts to set the pointer to beyond the file size, you must set the seek pointer to the end of the file. The

offset location of the seek pointer in the file is returned from the call to seek.

- **int close(int fd);** Closes the file corresponding to *fd*, commits all file transactions on this file, and unregisters *fd* from the user file descriptor table of the calling thread's TCB. The return value is 0 in success, otherwise -1.
- **int delete(String fileName);** Deletes the file specified by *fileName*. All blocks used by file are freed. If the file is currently open, it is not deleted and the operation returns a -1. If successfully deleted a 0 is returned.
- **int fsize(int fd);** Returns the size in bytes of the file indicated by *fd*.

Unless specified otherwise, each of the above system calls returns -1 when detecting an error.

5. Implementation

Superblock

The first disk block, block 0, is called the *superblock*. It is used to describe

1. The number of disk blocks.
2. The number of *inodes*.
3. The block number of the head block of the free list.

It is the OS-managed block. No other information must be recorded in and no user threads must be able to get access to the *superblock*.

```
class Superblock {
    public int totalBlocks; // the number of disk blocks
    public int totalInodes; // the number of inodes
    public int freeList;    // the block number of the free list's head

    public SuperBlock( int diskSize ) {
        ...
    }
    ...
}
```

Inodes

Starting from the blocks after the *superblock* will be the *inode* blocks. Each *inode* describes one file. Our *inode* is a simplified version of the Unix *inode*. It includes 12 pointers of the index block. The first 11 of these pointers point to direct blocks. The last pointer points to an indirect block. In addition, each *inode* must include (1) the length of the corresponding file, (2) the number of file (structure) table entries that point to this inode, and (3) the flag to indicate if it is unused (= 0), used(= 1), or in some other status (= 2, 3, 4, ...). 16 *inodes* can be stored in one block.

```
public class Inode {
    private final static int iNodeSize = 32;        // fix to 32 bytes
    private final static int directSize = 11;       // # direct pointers
```

```

public int length;                // file size in bytes
public short count;               // # file-table entries pointing to this
public short flag;                // 0 = unused, 1 = used, ...
public short direct[] = new short[directSize]; // direct pointers
public short indirect;            // a indirect pointer

Inode( ) {                        // a default constructor
    length = 0;
    count = 0;
    flag = 1;
    for ( int i = 0; i < directSize; i++ )
        direct[i] = -1;
    indirect = -1;
}

Inode( short iNumber ) {          // retrieving inode from disk
    // design it by yourself.
}

int toDisk( short iNumber ) {     // save to disk as the i-th inode
    // design it by yourself.
}
}

```

You will need a constructor that retrieves an existing *inode* from the disk into the memory. Given an *inode* number, termed *inumber*, this constructor reads the corresponding disk block, locates the corresponding *inode* information in that block, and initializes a new *inode* with this information.

The system must avoid any *inode* inconsistency among different user threads. There are two solutions to maintain the *inode* consistency:

1. Before an *inode* in memory is updated, check the corresponding *inode* on disk, read it from the disk if the disk has been updated by another thread. Thereafter, you should write back its contents to disk immediately. Note that the *inode* data to be written back include *int length*, *short count*, *short flag*, *short direct[11]*, and *short indirect*, thus requiring a space of 32 bytes in total. For this write-back operation, you will need the *toDisk* method that saves this *inode* information to the *iNumber*-th *inode* in the disk, where *iNumber* is given as an argument.
2. Create a **Vector<Inode>** object that maintains all *inode* on memory, is shared among all threads, and is exclusively access by each thread.

Root Directory

The "/" root directory maintains each file in a different directory entry that contains its file name (maximum 30 characters; 60 bytes in Java) and the corresponding *inode* number. The directory receives the maximum number of *inodes* to be created, (i.e., thus the max. number of files to be created) and keeps track of which *inode* numbers are in use. Since the directory itself is considered as a file, its contents are maintained by an *inode*, specifically *inode* 0. This can be located in the first 32 bytes of the disk block 1.

Upon booting ThreadOS, the file system instantiates the *Directory* class as the root directory through its constructor, reads the file from the disk that can be found through the *inode* 0 at 32 bytes of the disk block 1, and initializes the *Directory* instance with the file contents. Prior to shutdown, the file system must write back the *Directory* information onto the disk. The methods *bytes2directory()* and *directory2bytes* will initialize the *Directory* instance with a byte array read from the disk and converts the *Directory* instance into a byte array that will be thereafter written back to the disk.

```
public class Directory {
    private static int maxChars = 30; // max characters of each file name

    // Directory entries
    private int fsize[];           // each element stores a different file size.
    private char fnames[][];       // each element stores a different file name.

    public Directory( int maxInumber ) { // directory constructor
        fsizes = new int[maxInumber];    // maxInumber = max files
        for ( int i = 0; i < maxInumber; i++ )
            fsize[i] = 0;                // all file size initialized to 0
        fnames = new char[maxInumber][maxChars];
        String root = "/";               // entry(inode) 0 is "/"
        fsize[0] = root.length( );       // fsize[0] is the size of "/".
        root.getChars( 0, fsize[0], fnames[0], 0 ); // fnames[0] includes "/"
    }

    public void bytes2directory( byte data[] ) {
        // assumes data[] received directory information from disk
        // initializes the Directory instance with this data[]
    }

    public byte[] directory2bytes( ) {
        // converts and return Directory information into a plain byte array
        // this byte array will be written back to disk
        // note: only meaningful directory information should be converted
        // into bytes.
    }

    public short ialloc( String filename ) {
        // filename is the one of a file to be created.
        // allocates a new inode number for this filename
    }

    public boolean ifree( short iNumber ) {
        // deallocates this inumber (inode number)
        // the corresponding file will be deleted.
    }

    public short namei( String filename ) {
        // returns the inumber corresponding to this filename
    }
}
```

File (Structure) Table

The file system maintains the file (structure) table shared among all user threads. When a user thread opens a file, it follows the sequence listed below:

1. The user thread allocates a new entry of the user file descriptor table in its TCB. This entry number itself becomes a file descriptor number. The entry maintains a reference to a file (structure) table entry.
2. The user thread then requests the file system to allocate a new entry of the system-maintained file (structure) table. This entry includes the seek pointer of this file, a reference to the inode corresponding to the file, the inode number, the count to maintain #threads sharing this file (structure) table, and the access mode. The seek pointer is set to the front or the tail of this file depending on the file access mode.
3. The file system locates the corresponding *inode* and records it in this file (structure) table entry.
4. The user thread finally registers a reference to this file (structure) table entry in its file descriptor table entry of the TCB.

The file (structure) table entry should be:

```
public class FileTableEntry {           // Each table entry should have
    public int seekPtr;                 //   a file seek pointer
    public final Inode inode;           //   a reference to its inode
    public final short iNumber;         //   this inode number
    public int count;                   //   # threads sharing this entry
    public final String mode;           //   "r", "w", "w+", or "a"
    public FileTableEntry ( Inode i, short inumber, String m ) {
        seekPtr = 0;                   // the seek pointer is set to the file top
        inode = i;
        iNumber = inumber;
        count = 1;                     // at least on thread is using this entry
        mode = m;                       // once access mode is set, it never changes
        if ( mode.compareTo( "a" ) == 0 ) // if mode is append,
            seekPtr = inode.length;     // seekPtr points to the end of file
    }
}
```

The file (structure) table is defined as follows:

```
public class FileTable {

    private Vector table;               // the actual entity of this file table
    private Directory dir;              // the root directory

    public FileTable( Directory directory ) { // constructor
        table = new Vector( );         // instantiate a file (structure) table
        dir = directory;                // receive a reference to the Director
    }                                   // from the file system

    // major public methods
    public synchronized FileTableEntry falloc( String filename, String mode ) {
        // allocate a new file (structure) table entry for this file name
        // allocate/retrieve and register the corresponding inode using dir
    }
}
```

```

    // increment this inode's count
    // immediately write back this inode to the disk
    // return a reference to this file (structure) table entry
}

public synchronized boolean ffree( FileTableEntry e ) {
    // receive a file table entry reference
    // save the corresponding inode to the disk
    // free this file table entry.
    // return true if this file table entry found in my table
}

public synchronized boolean fempty( ) {
    return table.isEmpty(); // return if table is empty
}
// should be called before starting a format
}

```

File Descriptor Table

Each user thread maintains a user file descriptor table in its own TCB. Every time it opens a file, it allocates a new entry table including a reference to the corresponding file (structure) table entry. Whenever a thread spawns a new child thread, it passes a copy of its TCB to this child which thus has a copy of its parent's user file descriptor table. This in turn means the both the parent and the child refer to the same file (structure) table entries and eventually share the same files.

```

public class TCB {
    private Thread thread = null;
    private int tid = 0;
    private int pid = 0;
    private boolean terminate = false;

    // User file descriptor table:
    // each entry pointing to a file (structure) table entry
    public FileTableEntry[] ftEnt = null;

    public TCB( Thread newThread, int myTid, int parentTid ) {
        thread = newThread;
        tid = myTid;
        pid = parentTid;
        terminated = false;

        // The following code is added for the file system
        ftEnt = new FileTableEntry[32];
        for ( int i = 0; i < 32; i++ )
            ftEnt[i] = null; // all entries initialized to null
        // fd[0], fd[1], and fd[2] are kept null.
    }
}

```

Note that the file descriptors 0, 1, and 2 are reserved for the standard input, output, and error, and thus those entries must remain null.

6. Statement of Work

Design and implement the file system in *ThreadOS* as specified above. Run *Test5.java* which is found in the *ThreadOS* directory. This test program will test the functionality of every feature and system call listed above and check the data consistency of your file system. Some of the things that are tested include:

1. Does the filesystem allow one to read the same data from a file that was previously written?
2. Are malicious operations handled appropriately and proper error codes returned? Those operations include file accesses with a negative seek pointer, too many files to be opened at a time per a thread and over the system, etc.
3. Does the file system synchronize all data on memory with disk before a shutdown and reload those data from the disk upon the next boot?

7. What to Turn In

Softcopy (file uploads):

- Please turn in one report per collaboration group:
 1. All source code you have modified and created
 2. Results when you run your own test program
 3. Your file system specification including your assumptions and limitations
 4. Descriptions on internal design including inode, file table, etc.
 5. Consideration on performance estimation, current functionality, and possible extended functionality

Grade Guide

Assignment5 grade guide:

1. Report 30% (12 pts)
 - Result(s) when testing format/open/read/write/seek/close/fsize/del...4
 - Specification including your assumptions and limitations.....4
 - Descriptions on internal design including inode, file table, etc.....4
2. Correctness 70% (28 pts)
 - Completed Inode coding.....2
 - Completed Directory coding.....2
 - Completed File (structure) table coding.....2
 - Completed TCB user file descriptor table coding.....2
 - Sufficient coding of file system.....2

When running the instructor's Test5.java

 - Correct behavior of format.....2
 - Correct behavior of open.....1
 - Correct behavior of writing a few bytes (appx.16 bytes).....1
 - Correct behavior of close.....1
 - Correct behavior of reading a few bytes.....1

Correct behavior of appending a few bytes (appx.32 bytes).....	1
Correct behavior of seeking in a small file (appx.48 bytes).....	1
Correct behavior of reading/writing a small file (appx.48 bytes)...	0.5
Correct behavior of writing a lot of bytes (appx. 512 * 13 bytes)..	0.5
Correct behavior of reading a lot of bytes.....	0.5
Correct behavior of appending a few bytes to a large file.....	0.5
Correct behavior of seeking in a large file (appx. 512 * 13 + 48)..	0.5
Correct behavior of reading/writing a large file with wt.....	0.5
Correct behavior of delete.....	0.5
Correct behavior of creating over 40 files	0.5
Correct behavior of a parent and a child reading the same file.....	0.5
Correct behavior of reading the same file with two independent fds.	0.5
Other Correctness tests.....	5
Total.....	40

8. FAQ

This website could answer your questions. Please click [P5_FAQ](#)

(<https://canvas.uw.edu/courses/1309374/pages/p5-faq>) before emailing the professor :-).

9. Additional Materials

See [CSS430FinalProject.pdf](#) (<https://canvas.uw.edu/courses/1309374/files/54933429/download?wrap=1>)



and

[FinalProject.pptx](#) (<https://canvas.uw.edu/courses/1309374/files/56718694/download?wrap=1>)



(An alternative, possibly more clear, description of this assignment is at

<http://students.washington.edu/mwatson5/css-497/> (<http://students.washington.edu/mwatson5/css-497/>) but that version uses a C++ based ThreadOS, so you can read it for clarification, but do not rely on it.)